

MNIST Digit Classification with CNN

June 23, 2025

```
[4]: #MNIST Digit Classification with CNN (TensorFlow)
```

```
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import numpy as np
```

```
[5]: #Load MNIST dataset
```

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 0s

0s/step

```
[6]: # Preprocess the data
```

```
# Reshape to (28, 28, 1) and normalize pixel values to [0, 1]
```

```
x_train = x_train.reshape(-1, 28, 28, 1).astype("float32") / 255.0
```

```
x_test = x_test.reshape(-1, 28, 28, 1).astype("float32") / 255.0
```

```
[8]: #Build CNN model
```

```
model = models.Sequential([
    layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64, (3,3), activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax') # 10 digits (0-9)
])
```

```
[9]: #Compile the model
```

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
[10]: #Train the model
```

```
model.fit(x_train, y_train, epochs=5, validation_split=0.1)
```

```

Epoch 1/5
1688/1688          22s 12ms/step -
accuracy: 0.8947 - loss: 0.3469 - val_accuracy: 0.9855 - val_loss: 0.0477
Epoch 2/5
1688/1688          21s 12ms/step -
accuracy: 0.9836 - loss: 0.0518 - val_accuracy: 0.9892 - val_loss: 0.0379
Epoch 3/5
1688/1688          24s 14ms/step -
accuracy: 0.9903 - loss: 0.0316 - val_accuracy: 0.9898 - val_loss: 0.0362
Epoch 4/5
1688/1688          20s 12ms/step -
accuracy: 0.9923 - loss: 0.0247 - val_accuracy: 0.9872 - val_loss: 0.0433
Epoch 5/5
1688/1688          30s 18ms/step -
accuracy: 0.9945 - loss: 0.0171 - val_accuracy: 0.9927 - val_loss: 0.0327

```

[10]: <keras.src.callbacks.history.History at 0x7018b3fe63b0>

```

[11]: #Evaluate on test set
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy: {:.2f}%".format(test_acc * 100))

```

```

313/313          1s 4ms/step -
accuracy: 0.9895 - loss: 0.0357
Test Accuracy: 99.10%

```

```

[13]: #Visualize predictions on 5 sample test images
predictions = model.predict(x_test[:5])
predicted_labels = np.argmax(predictions, axis=1)
plt.figure(figsize=(10, 2))
for i in range(5):
    plt.subplot(1, 5, i+1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title(f"Pred: {predicted_labels[i]}")
    plt.axis('off')
plt.tight_layout()
plt.show()

```

1/1 0s 27ms/step



